

ASSIGNMENT 3: CNN

Rubén Martínez Sánchez de la Torre

Assignment Definition:

In this assignment, we will compare different Deep Convolutional Neural Networks in terms of:

- "number of parameters"
- "inference time"
- "performance".

You will be required to construct the following networks: 1. AlexNet 2. VGG-16 3. Your custom Deep CNN (CNN_5x5_Net): use only 5x5 convolutional kernels (hint: in VGG-16, all convolutional kernels were 3x3)

- Check the number of parameters vs inference time (by generating random noise image and feed-forward through the networks)

You can find the whole notebook attached with this document and in the following Jovian link and :

<https://jvn.io/rubensilver/dac35e1a7479468e90af9e5d114ca732>

AlexNet: Total params: 23,981,450 Trainable params: 23,981,450
Non-trainable params: 0

VGG-16: Total params: 134,301,514 Trainable params: 134,301,514
Non-trainable params: 0

CNN_5x5_Net: Total params: 160,453,450 Trainable params:
160,453,450 Non-trainable params: 0

Average Inference time for AlexNet: 0:00:00.157790

Average Inference time for VGG-16: 0:00:00.932858

Average Inference time for CNN_5x5_NET 0:00:01.201763

- Explain the results: why one was faster than the another, what was the key point?

In the previous section we can see that AlexNet is the network with the smallest number of parameters and therefore with the fastest inference time. The VGG-16 network has more parameters than AlexNet but less than CNN_5x5_Net and we can observe that its inference time is greater than the inference time of Alexnet but less than the inference time of CNN_5x5_Net.

The key point was the number of trainable parameters of each network.

The formula to get the size of the output feature map after applying a convolutional kernel over an input image / input feature map is:

$$FMsize = \frac{W - K + 2P}{S} + 1 \quad \text{where } W \text{ is the size of weight or the height of the input}$$

image or feature map, K is the kernel size that will be applied over the image, P is the padding and S is the strides.

For example, if you have an input image of $12 \times 12 \times 3$ and you're going to apply a kernel size $5 \times 5 \times 3$ using padding = 0 and strides = 1, the size of the output feature map will be:

$$FMsize = \frac{12 - 5 + 2 \times 0}{1} + 1 = 8$$

Using the formula above the original $12 \times 12 \times 3$ image will become a $8 \times 8 \times 1$ image.

The number of multiplications needed will be $5 \times 5 \times 3 = 75$ multiplications every time the kernel moves and as we have to move the kernel 8×8 times the final number of multiplications will be $5 \times 5 \times 3 \times 8 \times 8 = 4800$.

Each kernel is going to learn to detect a visual feature and if we want to train, for example, 64 kernels of 5×5 over an image of 3 channels then the number of weights will be: $32 * 5 * 5 * 3 = 2400$ and the memory needed to store the output feature map will be $8 * 8 * 32 = 2048$. We could multiply this value by 4 to get the raw number of bytes because each floating point number is composed by 4 bytes or by 8 if we are using double precision.

If we set $S = 1$ and we pad the input image / input feature map with 0's so that

$$P = \frac{(K - 1)}{2} \quad \text{then we are going to preserve the input size.}$$

If we use smaller kernel sizes we will have smaller receptive fields and we will capture smaller and complex features in the images.

If we use wider convolutions, I mean, layers with more kernels it will be easier for the GPU to process a single big chunk of data instead of a lot of smaller ones (that would be the case of having more layers with fewer kernels per layer).

- Add "Batch Normalization" and/or "Dropout" layers to the networks (AlexNet, VGG-16, Your custom CNN_5x5_Net)

This code can be found in the attached notebook.

When we train a Deep Neural Network we usually normalize the input data so that it follows a normal distribution with zero mean and unitary variance to assure that all input data is in the same range and to prevent early saturation in non linear function among other benefits.

But as the distribution of the activations is going to be hanging during the training process we are going to have a problem in the intermediate layers. Each layer is going

to learn to adapt itself to a new distribution each train step. This situation is known as **internal covariate shift** and it will slow down the training process.

To address that problem we are going to use a **Batch Normalization layer** which is a technique to normalize the input of each layer to prevent the internal covariate shift.

This layer should be applied before of the non linearity of the current layer according to the original paper:

"We add the BN transform immediately before the non-linearity, by normalizing $x = Wu+b$. We could have also normalized the layer inputs u , but since u is likely the output of another nonlinearity, the shape of its distribution is likely to change during training, and constraining its first and second moments would not eliminate the covariate shift. In contrast, $Wu + b$ is more likely to have a symmetric, non-sparse distribution, that is "more Gaussian"; normalizing it is likely to produce activations with a stable distribution."

The process to compute the output of this layer is:

- * Compute the mean and variance of the inputs of the layer where BN is being applied:

$$\text{Batch mean} \rightarrow \mu = \frac{1}{n} \cdot \sum_{i=1}^n x_i$$

$$\text{Batch variance} \rightarrow \sigma^2 = \frac{1}{n} \cdot \sum_{i=1}^n (x_i - \mu)^2$$

- * Normalize the inputs of the layer:

$$\bar{x}_i = \frac{x_i - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

- * Get the output of the BN layer scaling and shifting the normalized inputs:

$$y_i = \gamma \cdot \bar{x}_i + \beta \quad \text{where } \gamma \text{ and } \beta \text{ are learned during training.}$$

When we use a **Dropout layer**, at each training step, each neuron will have a probability $1 - p$ to be removed in that step or a probability p to be kept. The main benefit of this method is to prevent overfitting.

I have found a paper whose link is <https://arxiv.org/pdf/1905.05928.pdf> that proposes a layer composed by a BN layer followed by a Dropout layer to speed up the training process:

from tensorflow.keras.layers import BatchNormalization, Dropout

def Independent_Components_layer(inputs, p):

 x = BatchNormalization(inputs)

 x = Dropout(p)(x)

 return x

- Check how does inference time is changing after adding those layers Note:
Framework: we will be using Keras in this assignment.

AlexNet: Total params: 24,019,722 Trainable params: 24,000,586
Non-trainable params: 19,136

VGG-16: Total params: 134,351,178 Trainable params:
134,326,346 Non-trainable params: 24,832

CNN_5x5_Net: Total params: 160,503,114 Trainable params:
160,478,282 Non-trainable params: 24,832

Average Inference time for AlexNet: 0:00:00.167202

Average Inference time for VGG-16: 0:00:01.012128

Average Inference time for CNN_5x5_NET 0:00:01.301401